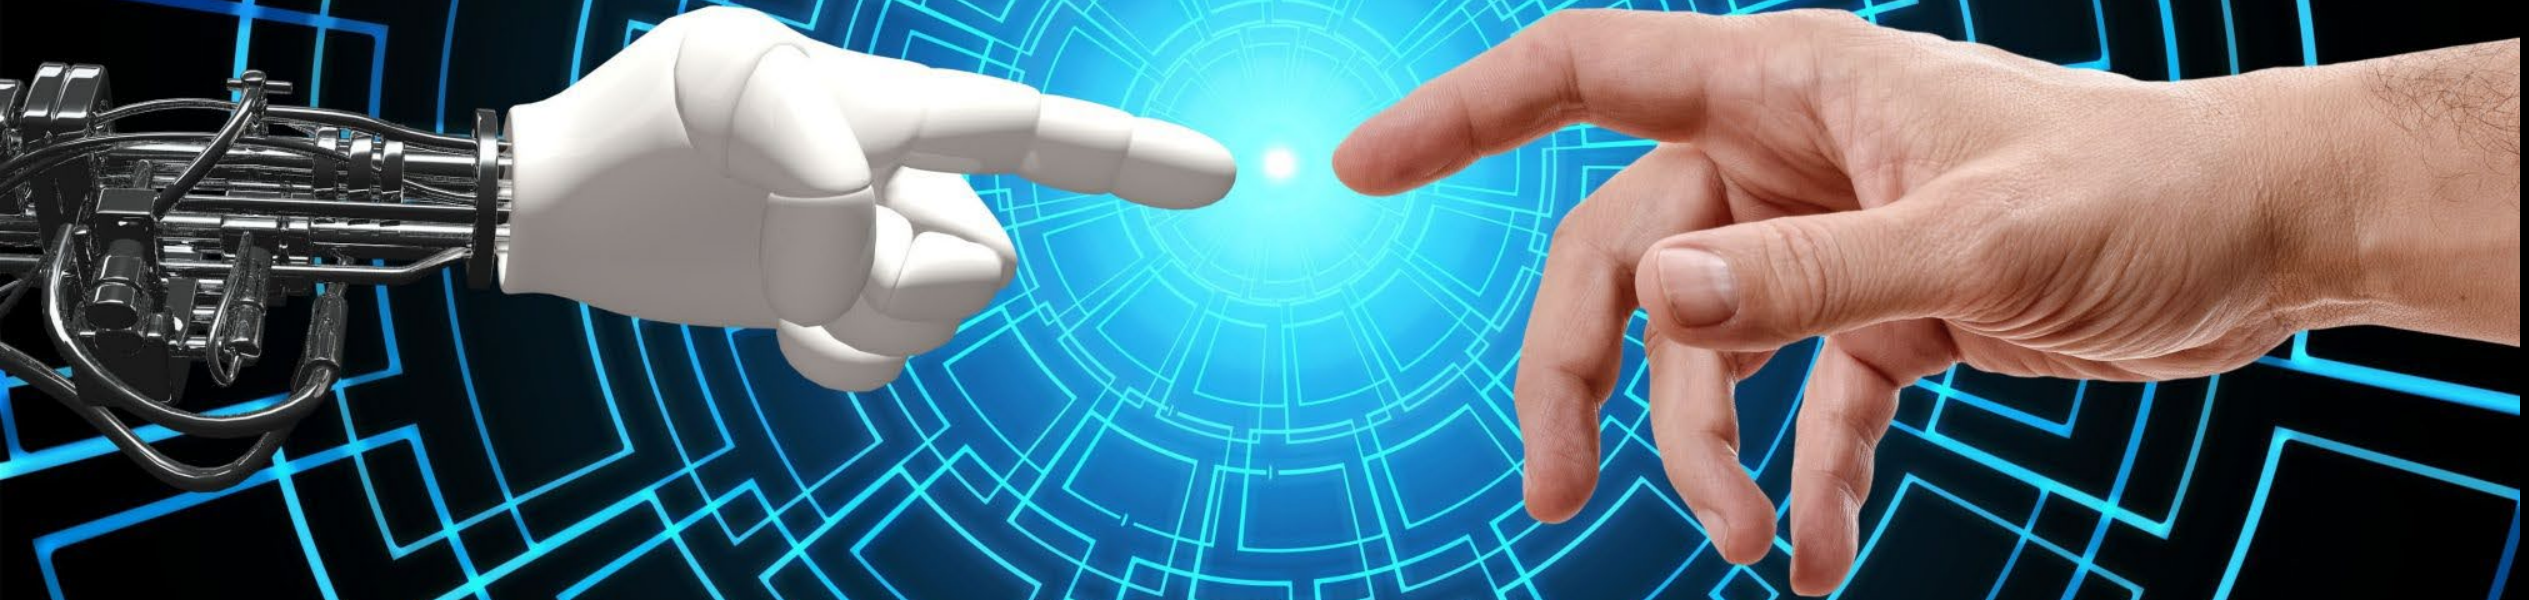


“Develop a passion for learning.”

3. MLOps for Containers and Edge Devices



Machine Learning for Operations

Module Topics

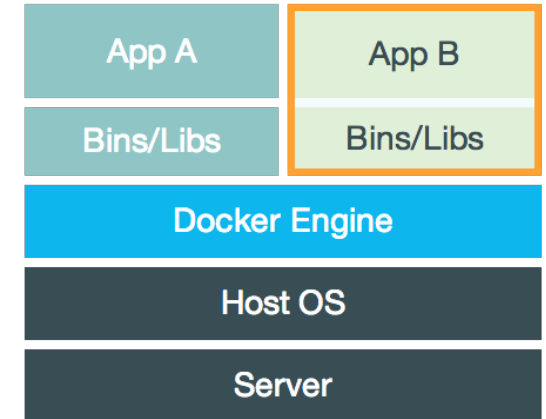
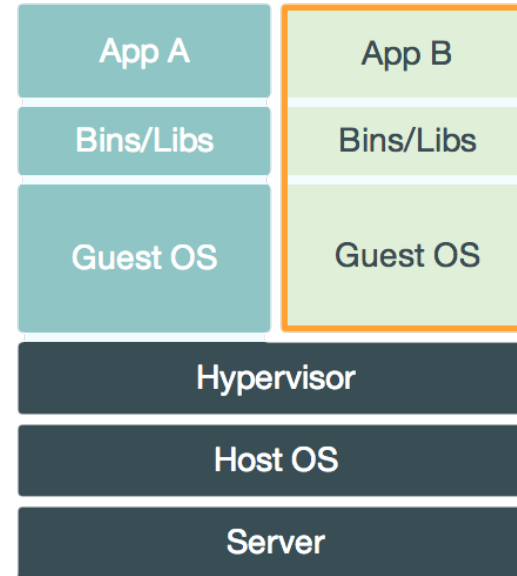
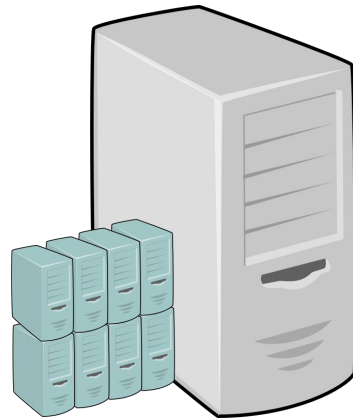
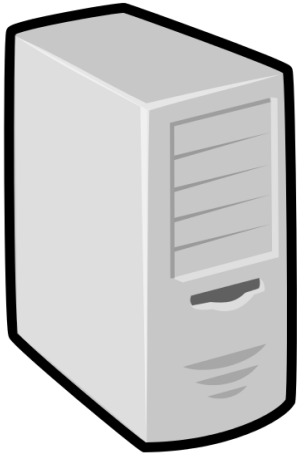
- Containers
- Serving a trained model over HTTP
- Edge Devices
- Coral
- AWS edge computing
- Tensor Flow Hub



Containers

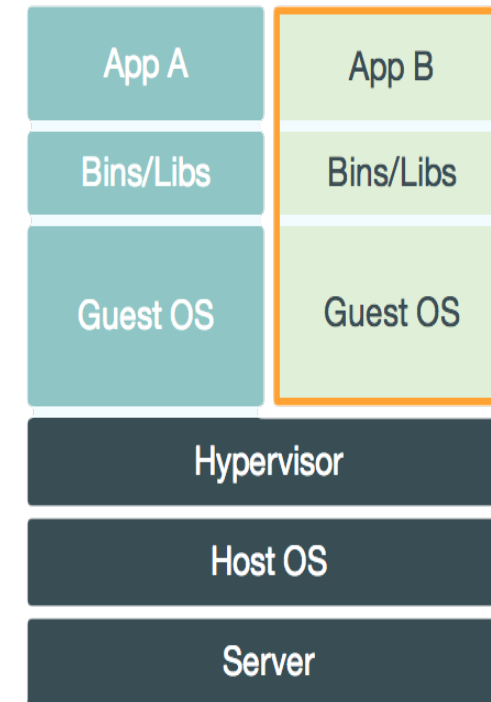


Evolution of the Data Center



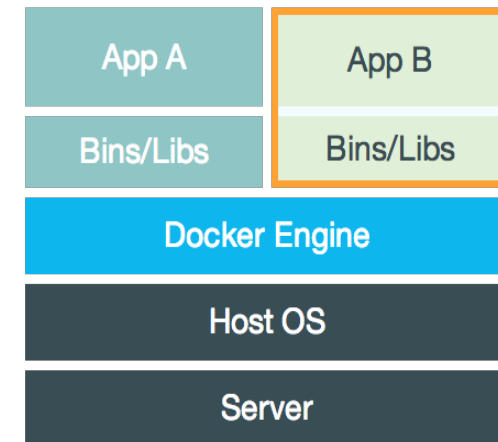
Virtual Machines

- Works in layers:
 - Hypervisor : that allows virtualization
 - Virtual hardware: simulates hardware
 - Guest OS: Linux, Windows etc.
- Performance penalty for keeping a full copy of the OS
 - Does provide complete isolation for each VM
 - Can mix OSes on the same physical host



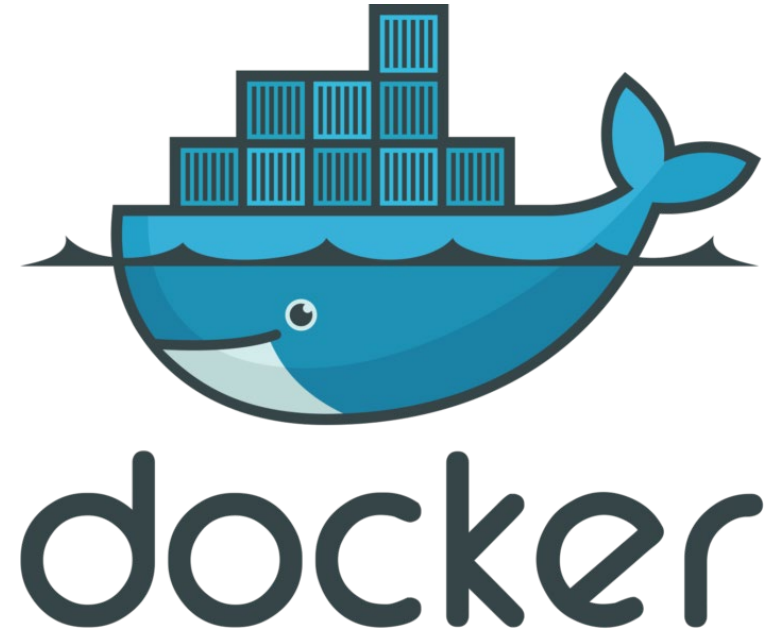
Container

- Virtualize the OS
 - Like how hypervisors virtualizes hardware
 - Shares the kernel with the Host OS
- Lightweight package that runs as a process on the host
 - Self-contained
 - Has everything needed to run
- Docker can run Linux containers on Windows
 - Uses a minimal Linux VM (LinuxKit)
 - Now uses WLS (Windows Linux Subsystem)



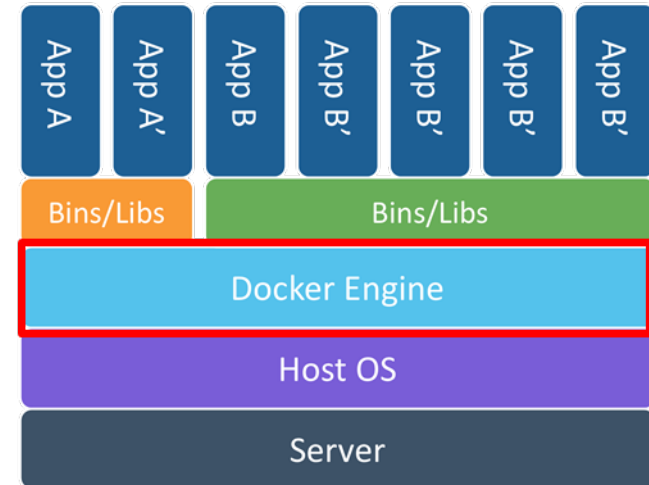
Docker

- Docker was introduced in 2015
- Built on Linux containers
- But had an ecosystem that allowed easy use and deployment
- Has become the de facto container standard



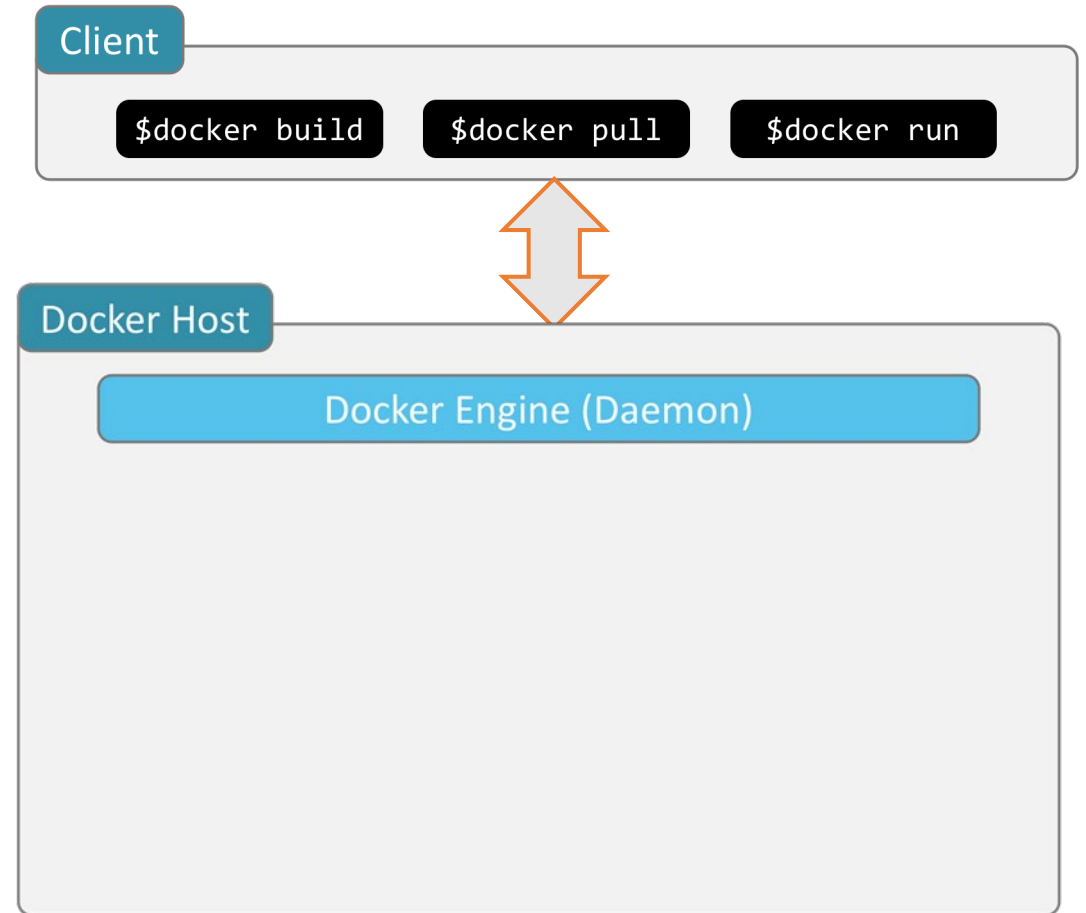
Containers: Docker Engine

- Lightweight runtime program to build, ship, and run Docker containers
 - Also known as Docker Daemon
 - Uses Linux Kernel namespaces and control groups
 - Linux Kernel (≥ 3.10)
 - Namespaces provide an isolated workspace



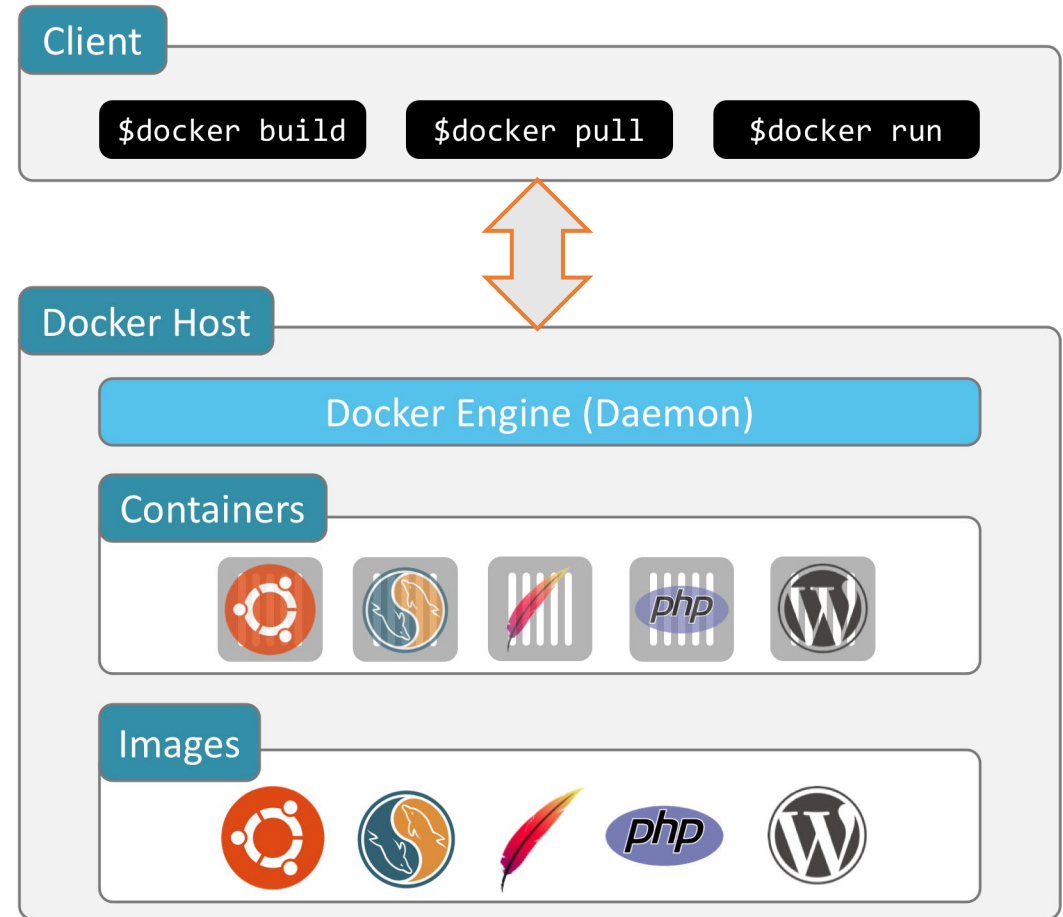
Docker Client and Engine

- The Docker Client is the docker binary
 - Primary interface to the Docker Host
 - Accepts commands and communicates with the Docker Engine (Daemon)



Containers: Docker Client

- Lives on a Docker host
- Creates and manages containers on the host



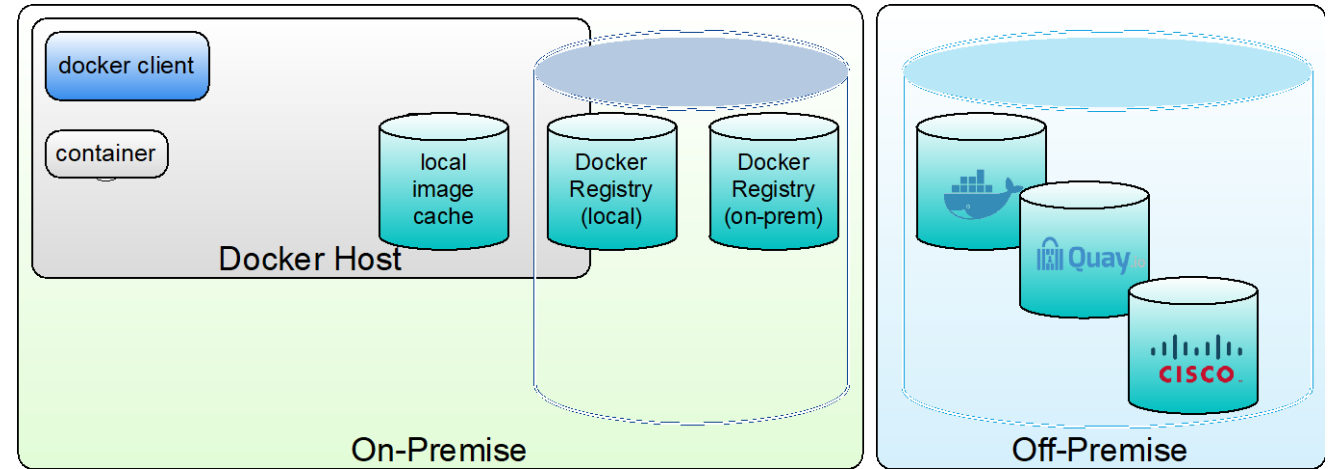
Docker Registries

- Image Storage & Retrieval System
- Docker Engine Pushes Images to a Registry
- Version Controlled
- Docker Engine Pulls Images to Run



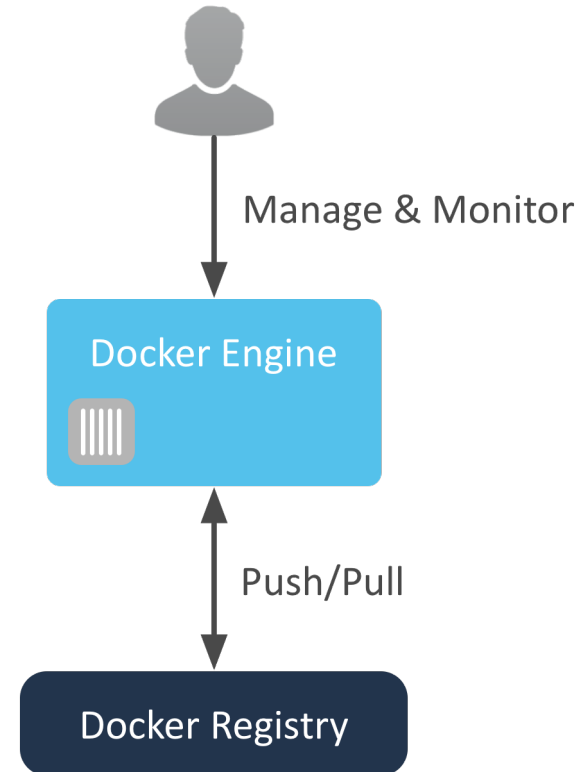
Docker Registries

- Types of Docker Registries
 - Local Docker Registry
 - On Docker Host
 - Remote Docker Registry
 - On-Premise/Off-Premise
 - Docker Hub
 - Off-Premise



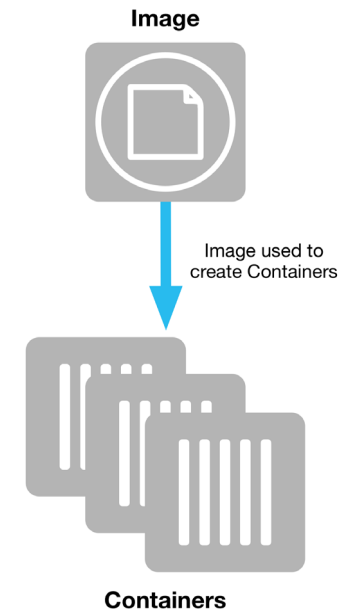
Pulling and Pushing

- Push
 - Upload an image to a registry
- Pull
 - Download an image from a registry
- Images are versioned in the registry



Docker Image Terminology

- Hierarchy of files, with meta-data for how to create/run a container
 - Read-only template used to create containers
 - Can be exported or modified to be used in new images
 - Created manually or through automated processes
 - Stored in a Registry (Docker Hub, Docker Trusted Registry, etc.)

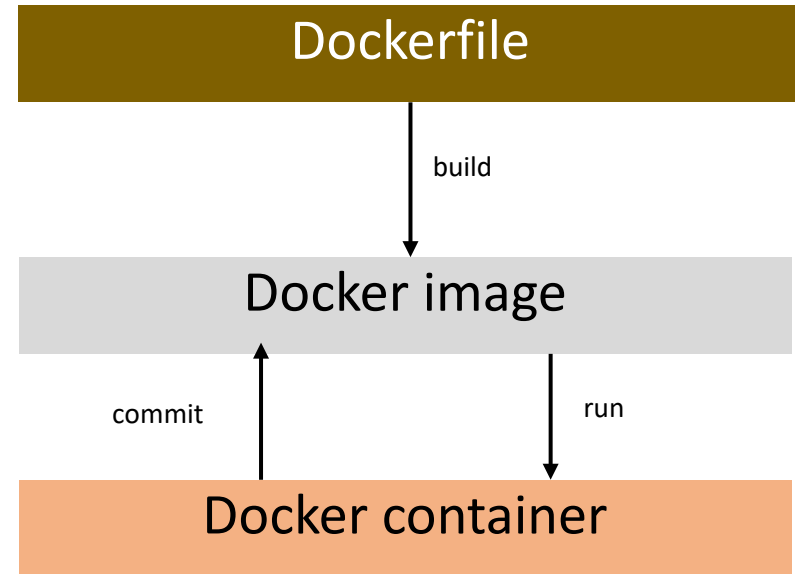


Inside a Docker Image/Container

- Docker images are built up in layers of file systems
 - Each layer is immutable
 - Each layer has a unique ID
 - The final image also has a unique ID
 - Immutable layers are shared/reused across images
- When a container is created from an image
 - An additional writeable layer is created
 - This maintains changes to the file system of a running container
 - All the immutable layers are shared across containers
- This architecture keeps containers light and small

Dockerfile

- Configuration File (script) for creating images. Defines:
 - Existing image to be the starting point
 - Set of instructions to augment that image (each of which results in a new layer of the file system)
 - Meta-data such as ports exposed
 - The command to execute when the image is run



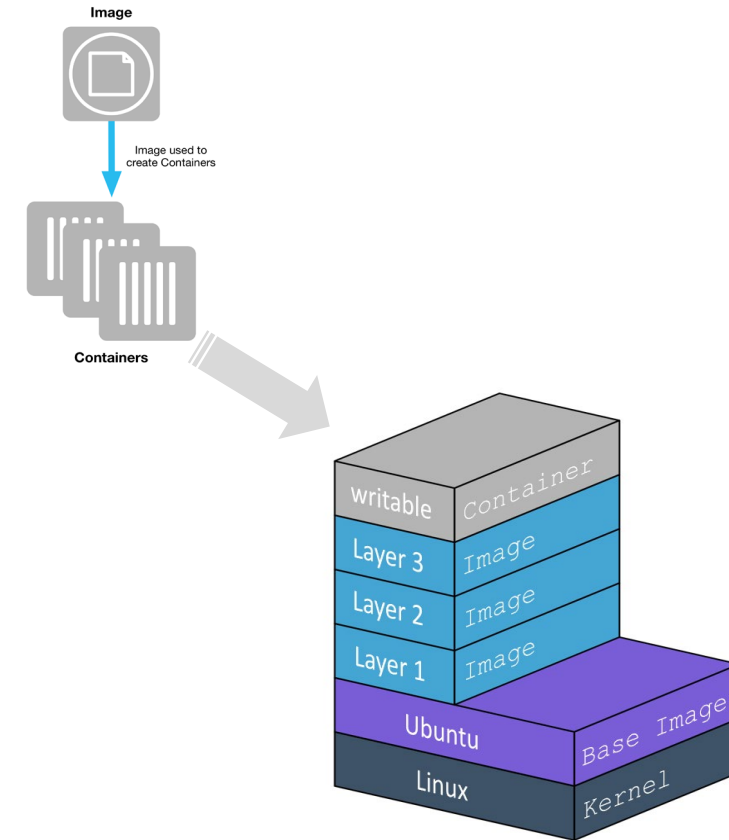
Deploying a Docker Application

Lab 4

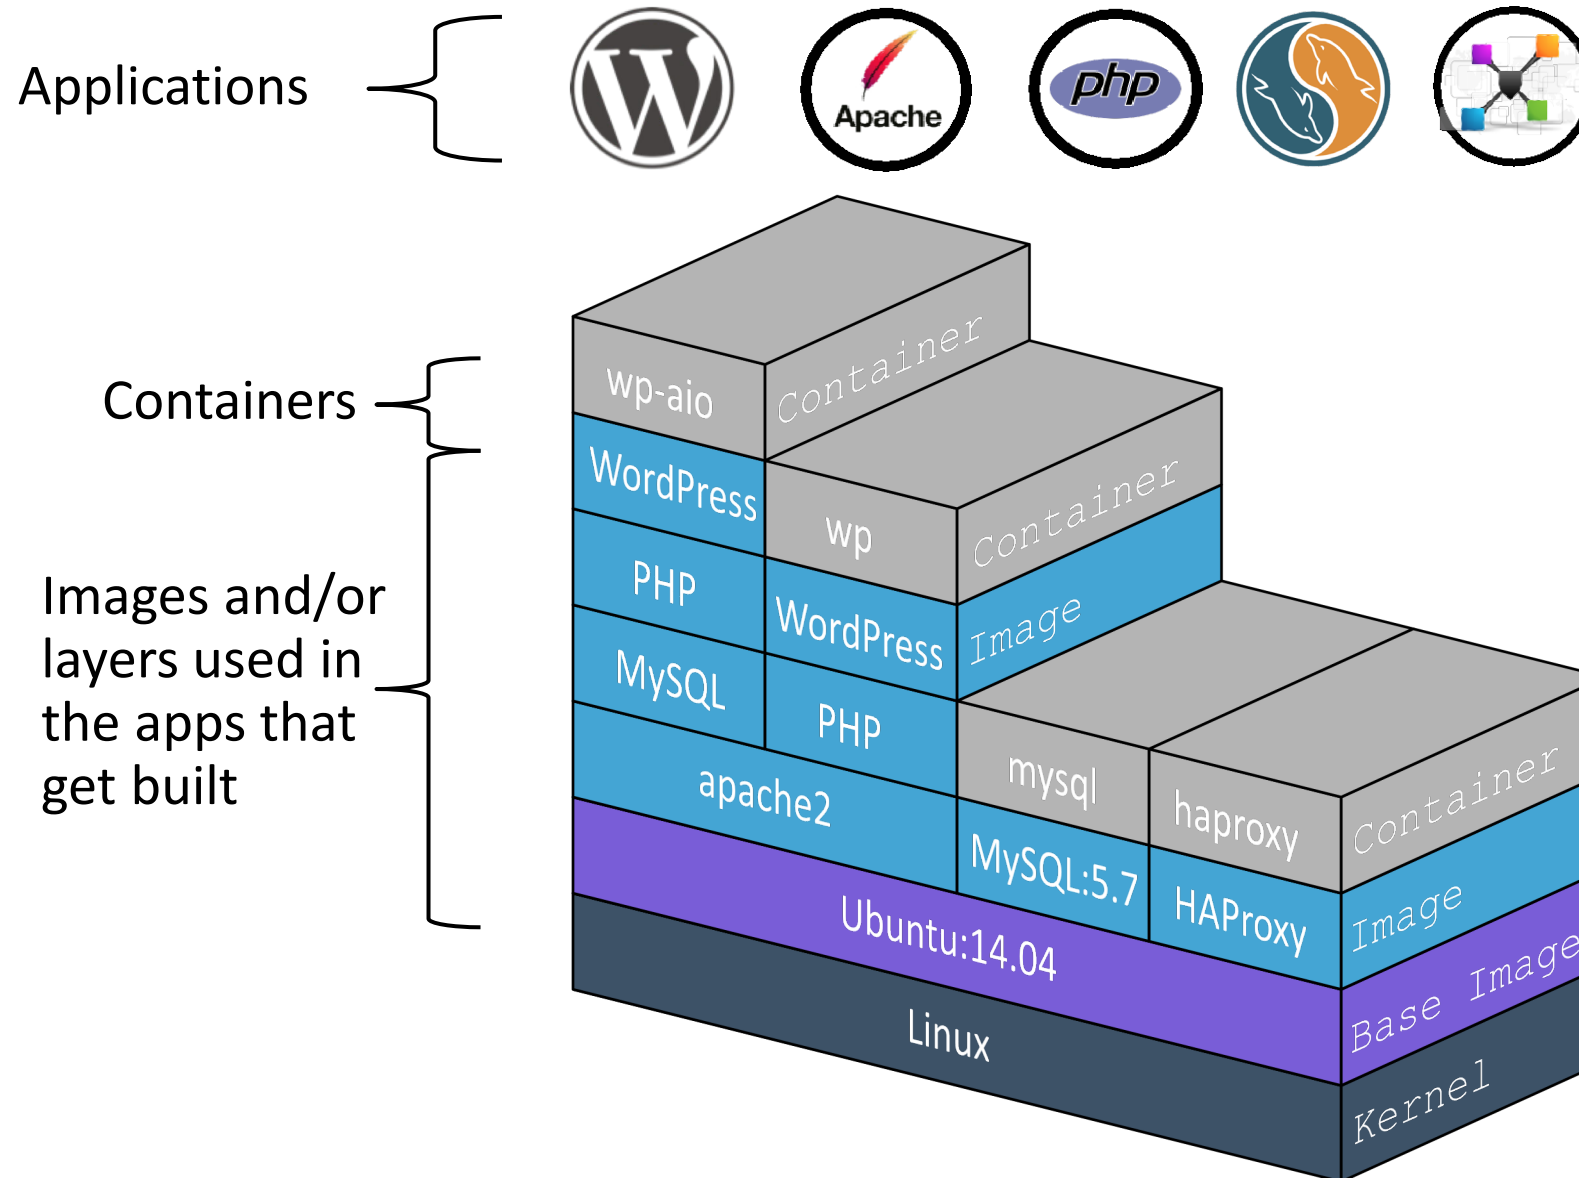


Container

- Runtime instance of an image plus a read/write layer

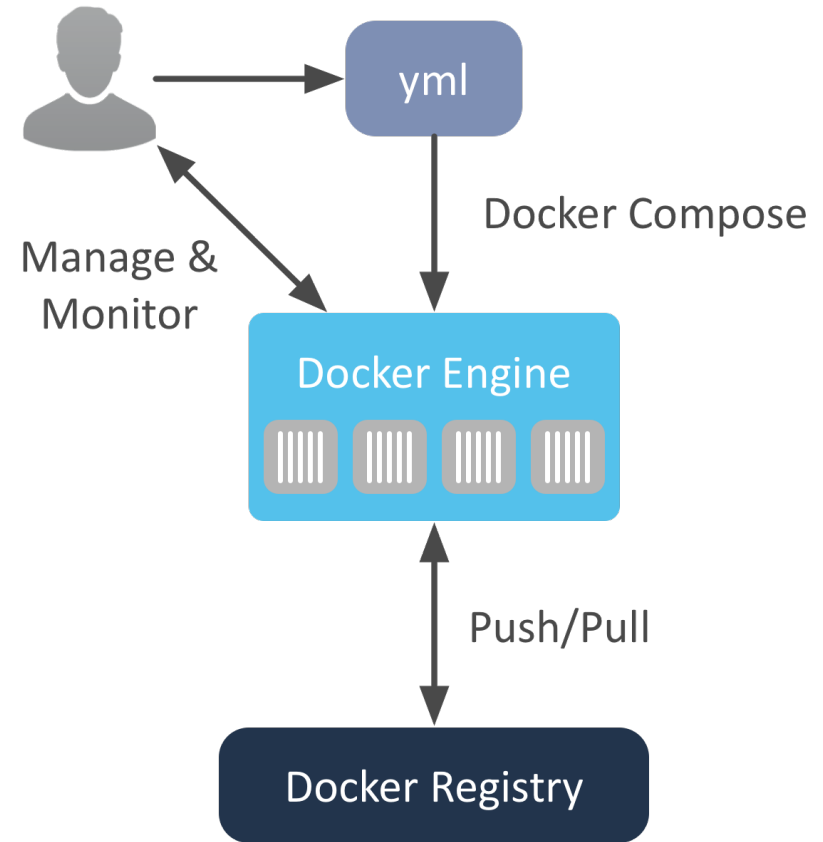


Inside a Docker Image/Container



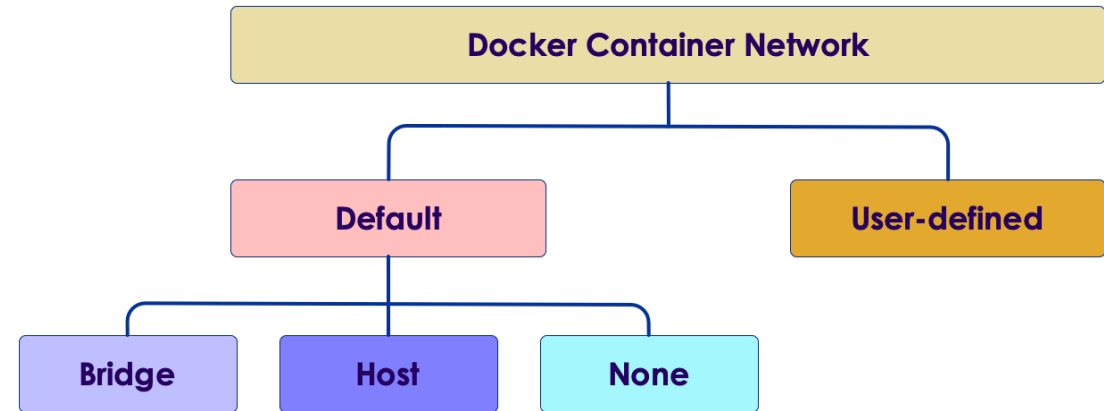
Docker Compose

- Tool to create and manage multi-container applications
- Applications defined in a single file: *docker-compose.yml*
- Transforms applications into individual containers that are linked together
- Compose will start all containers in a single command



Docker Networking

- Three default networks
- Bridge
 - Uses IPs in the range of 172.17.x.x
 - Local network for inter-container communication
- Host and None are not widely used
 - Use same network as host or none
- Additional networks can be defined



Port Mapping

- A container's functionality is accessed through a port
- Container ports on the bridge network are inaccessible
- Port mapping associates a container port with a host port
- This is used when a container is external interface
 - The following code maps the container port 80 to the host port 8998
 - Now the app can be accessed at the host port

```
docker run -d --name app -p 8998:80 app
```

```
curl localhost:8998
```

Docker Storage

- Containers have a writeable layer
 - This is ephemeral – it is lost when the container ends
- In order to have data persist, there are several options
 - Bind, where host directory is mounted in the container, like a NFS mount
 - The preferred mechanism is a Docker Volume
- Volumes are managed by Docker
 - Can use remote storage like AWS S3
 - Default location is `/var/lib/Docker`
 - Allows persistence and sharing of data between containers

Building a Docker Application

Lab 5

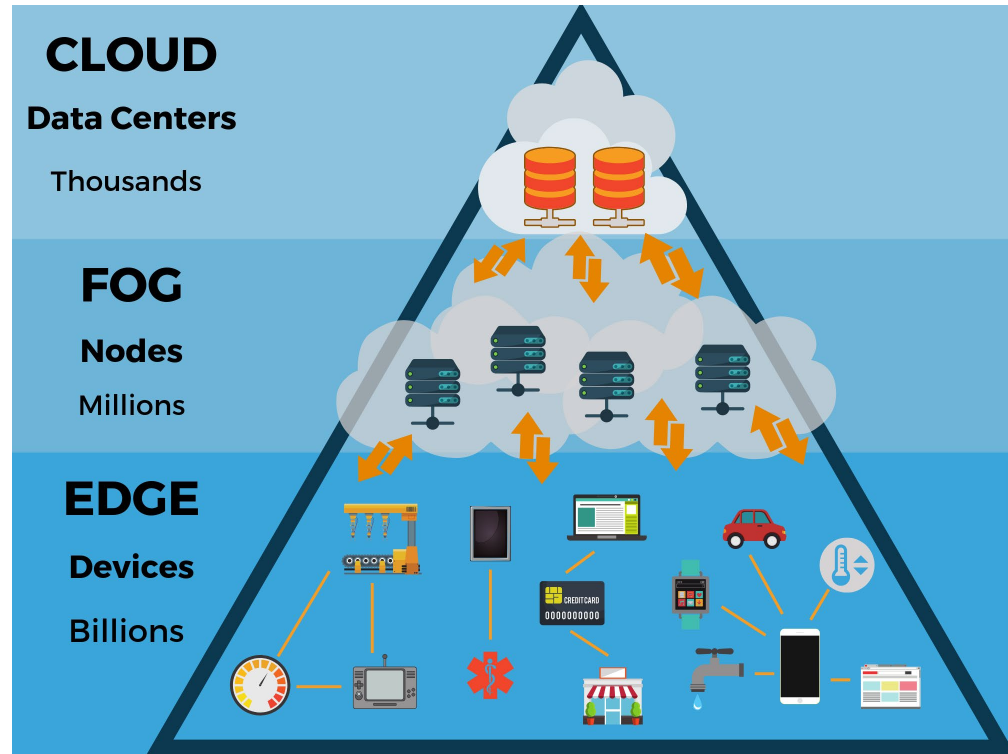


Edge Devices



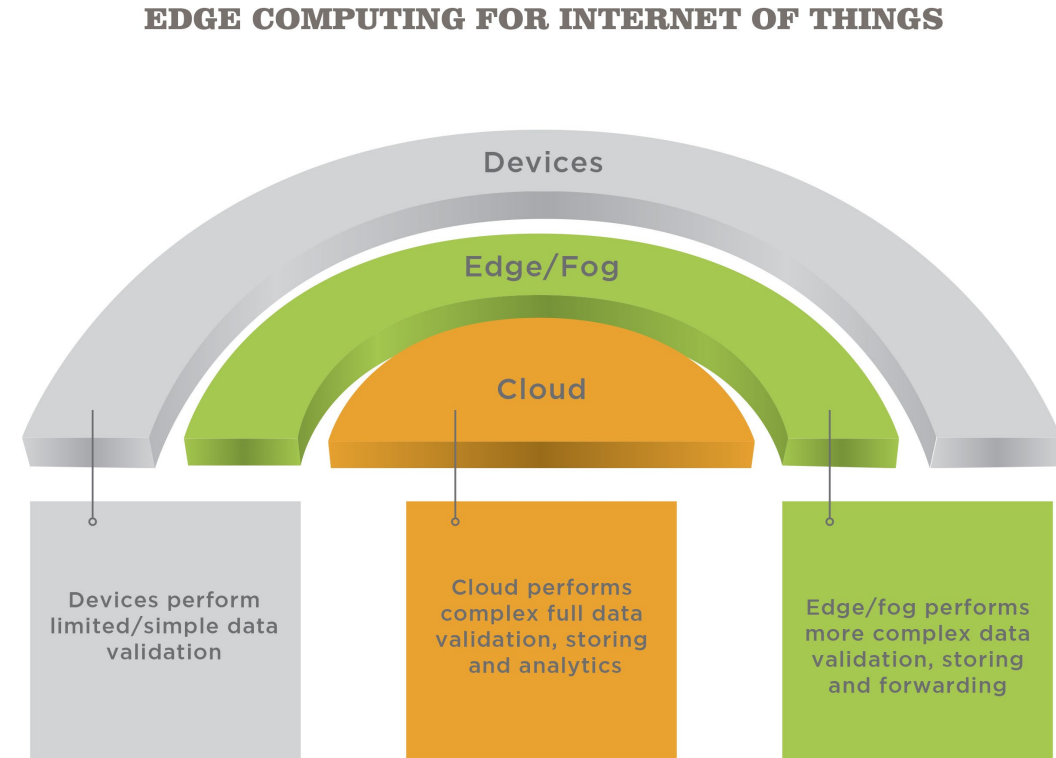
Edge Devices

- Devices with computing capability
 - Smart phones
 - Embedded processors
 - Sensors
 - Raspberry Pi
- Operate on the “edge” of cyberspace
 - Interface to the physical world
 - Cyber-physical connections
 - Internet of Things



Edge Computing


- Move the application closer to the user
 - Pushing the ML model out to edge devices
- Advantages
 - Reduces latency
 - Improves throughput
 - Parallel operations
 - Removes a point of failure
- Disadvantages
 - Relies on edge device capabilities and health
 - Updates and deployment can be problematic



The Coral Project

- Located at <https://coral.ai/>
- Focused on enabling privacy-preserving Edge ML with low-power, high performance products
- USB Accelerator
 - *“The Coral USB Accelerator adds an Edge TPU coprocessor to your system, enabling high-speed machine learning inferencing on a wide range of systems, simply by connecting it to a USB port.”*

The Coral Project




Dev Board

A single-board computer with a removable system-on-module (SoM) featuring the Edge TPU.

- Supported OS: Mendel Linux (derivative of Debian)
- Supported Framework: [TensorFlow Lite](#)
- Languages: Python and C++
- Works with [AutoML Vision Edge](#)
- [Datasheet](#)

→ View product




USB Accelerator

A USB accessory featuring the Edge TPU that brings ML inferencing to existing systems.

- Supported host OS: Debian Linux
- Compatible with Raspberry Pi boards
- Supported Framework: [TensorFlow Lite](#)
- Works with [AutoML Vision Edge](#)
- [Datasheet](#)

→ View product



Camera

5-megapixel camera module that's compatible with the Coral Dev Board.

- Omnivision Sensor
- Autofocus
- Auto exposure control
- 25 mm x 25 mm
- [Datasheet](#)

→ View product

Coral Project

Walkthrough and Deep Dive



exit

$$s(n-k)$$

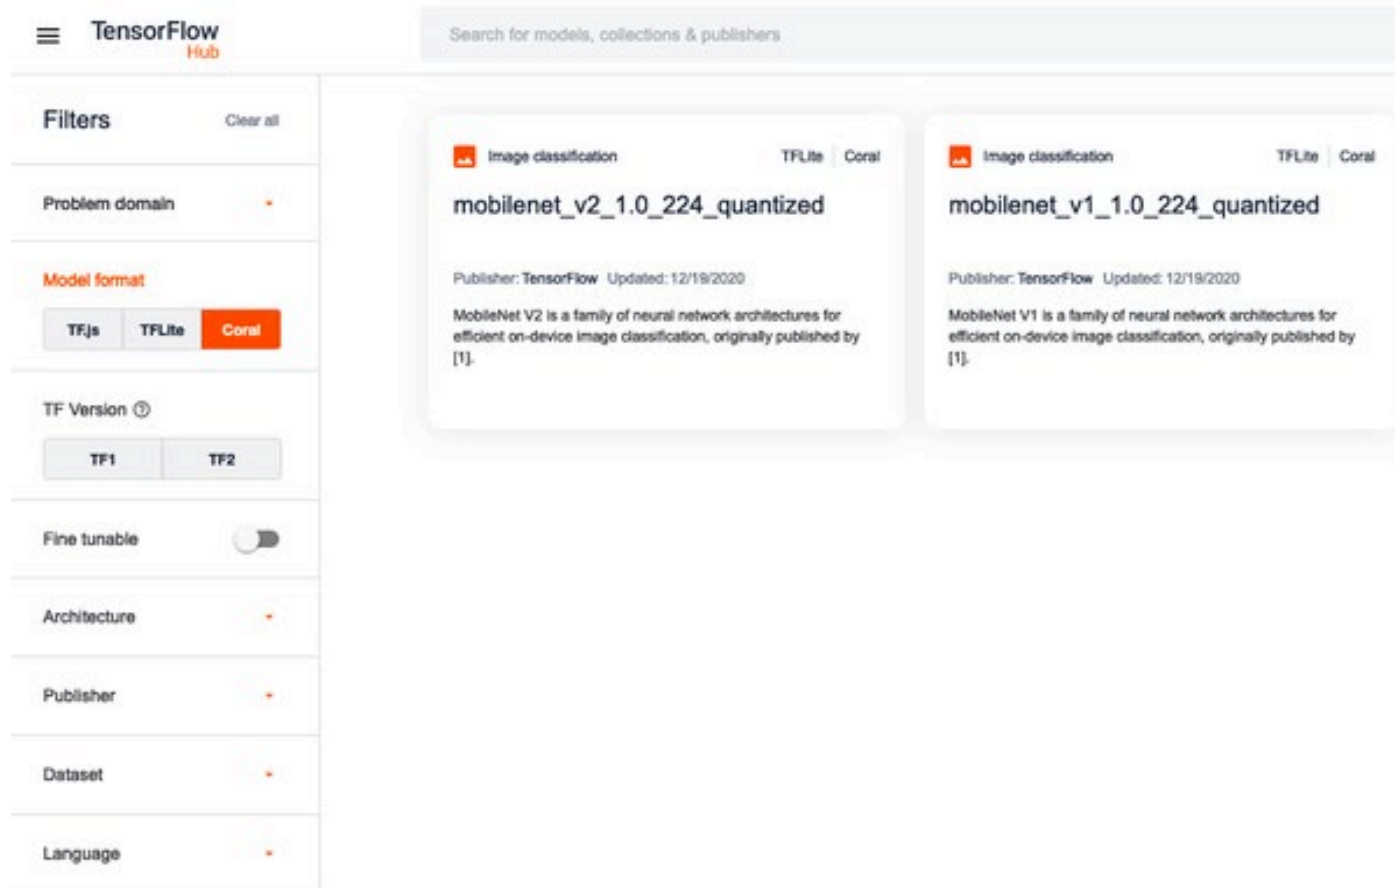
$$s(n-1)$$

$$s(n) = k_1 \cdot s(n-1) + k_2 \cdot$$

$$\cdot s(n-j) +$$

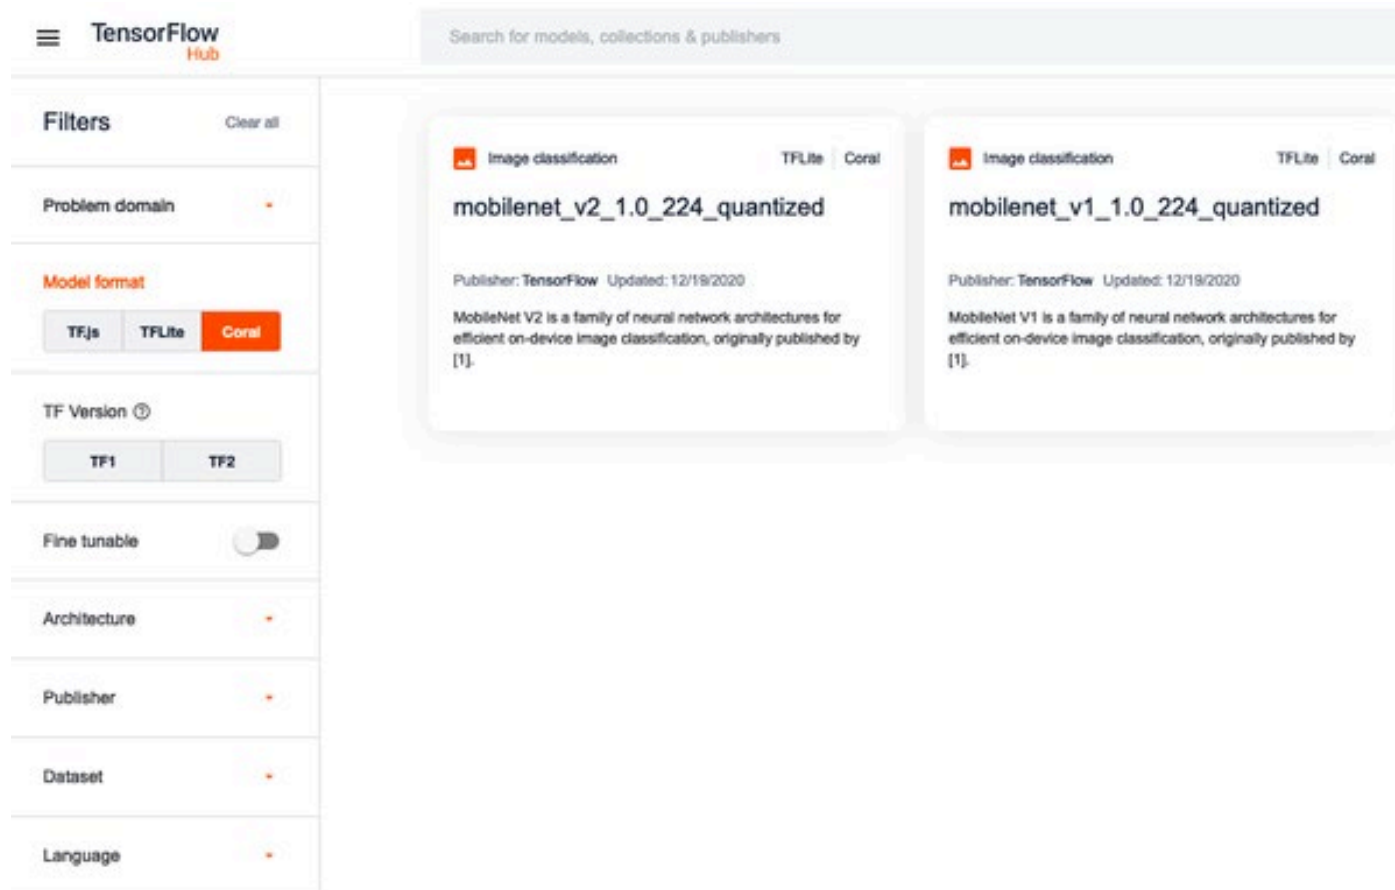
TFHub – TensorFlow Hub

- Located at <https://tfhub.dev/>
- Repository of thousands of pretrained models ready to be used.
 - For the Coral Edge TPU, not all models will work, though.

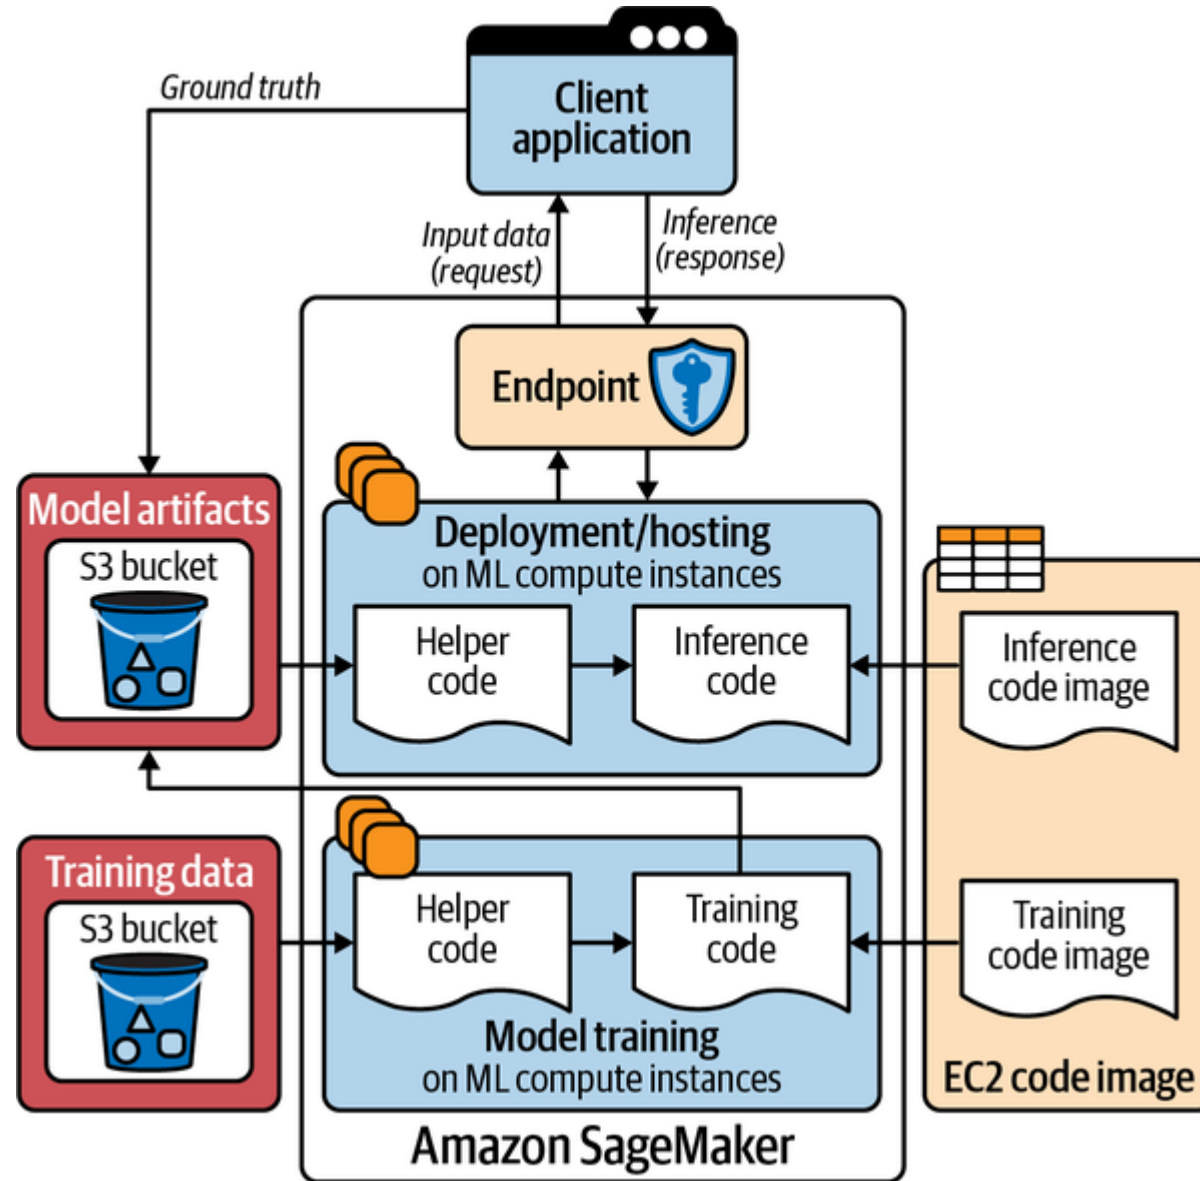


TFHub – TensorFlow Hub

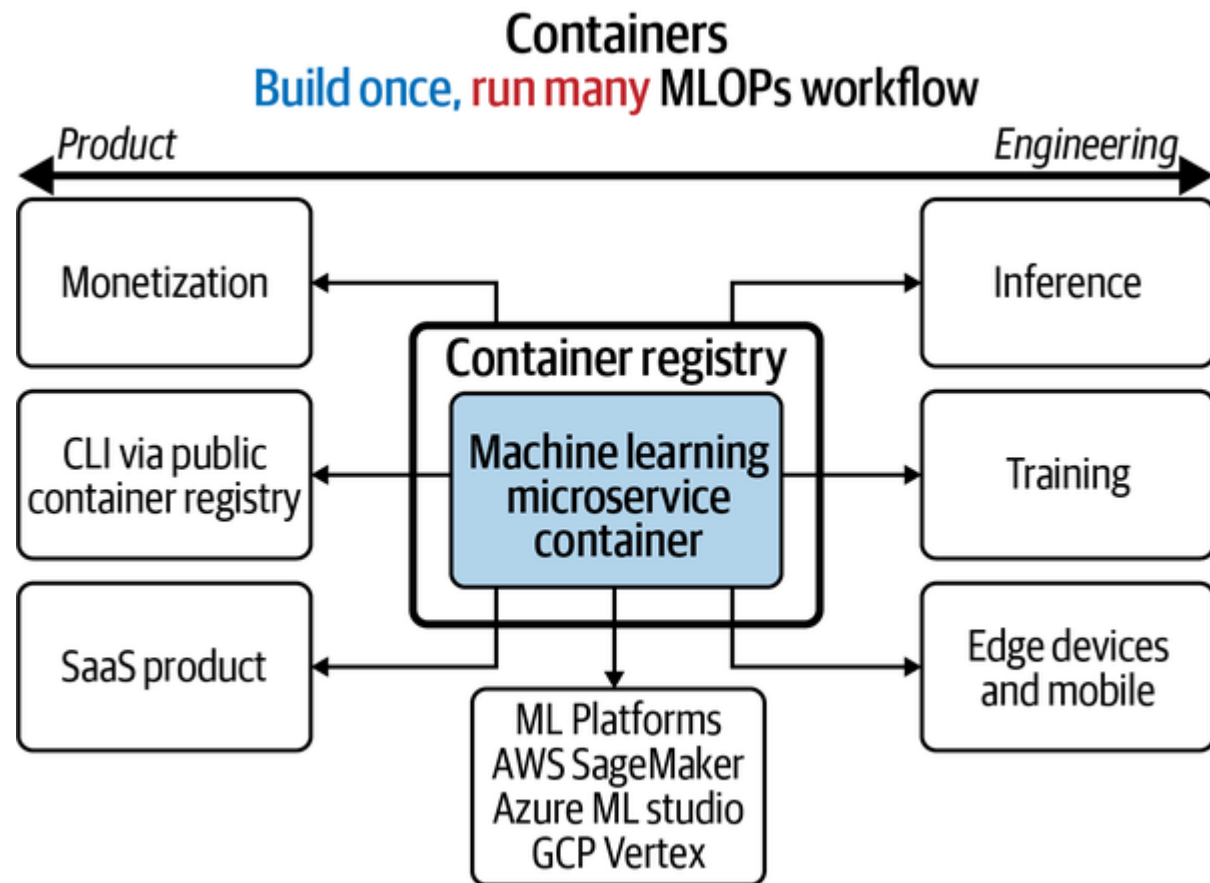
- Located at <https://tfhub.dev/>
- Repository of thousands of pretrained models ready to be used.
 - For the Coral Edge TPU, not all models will work, though.

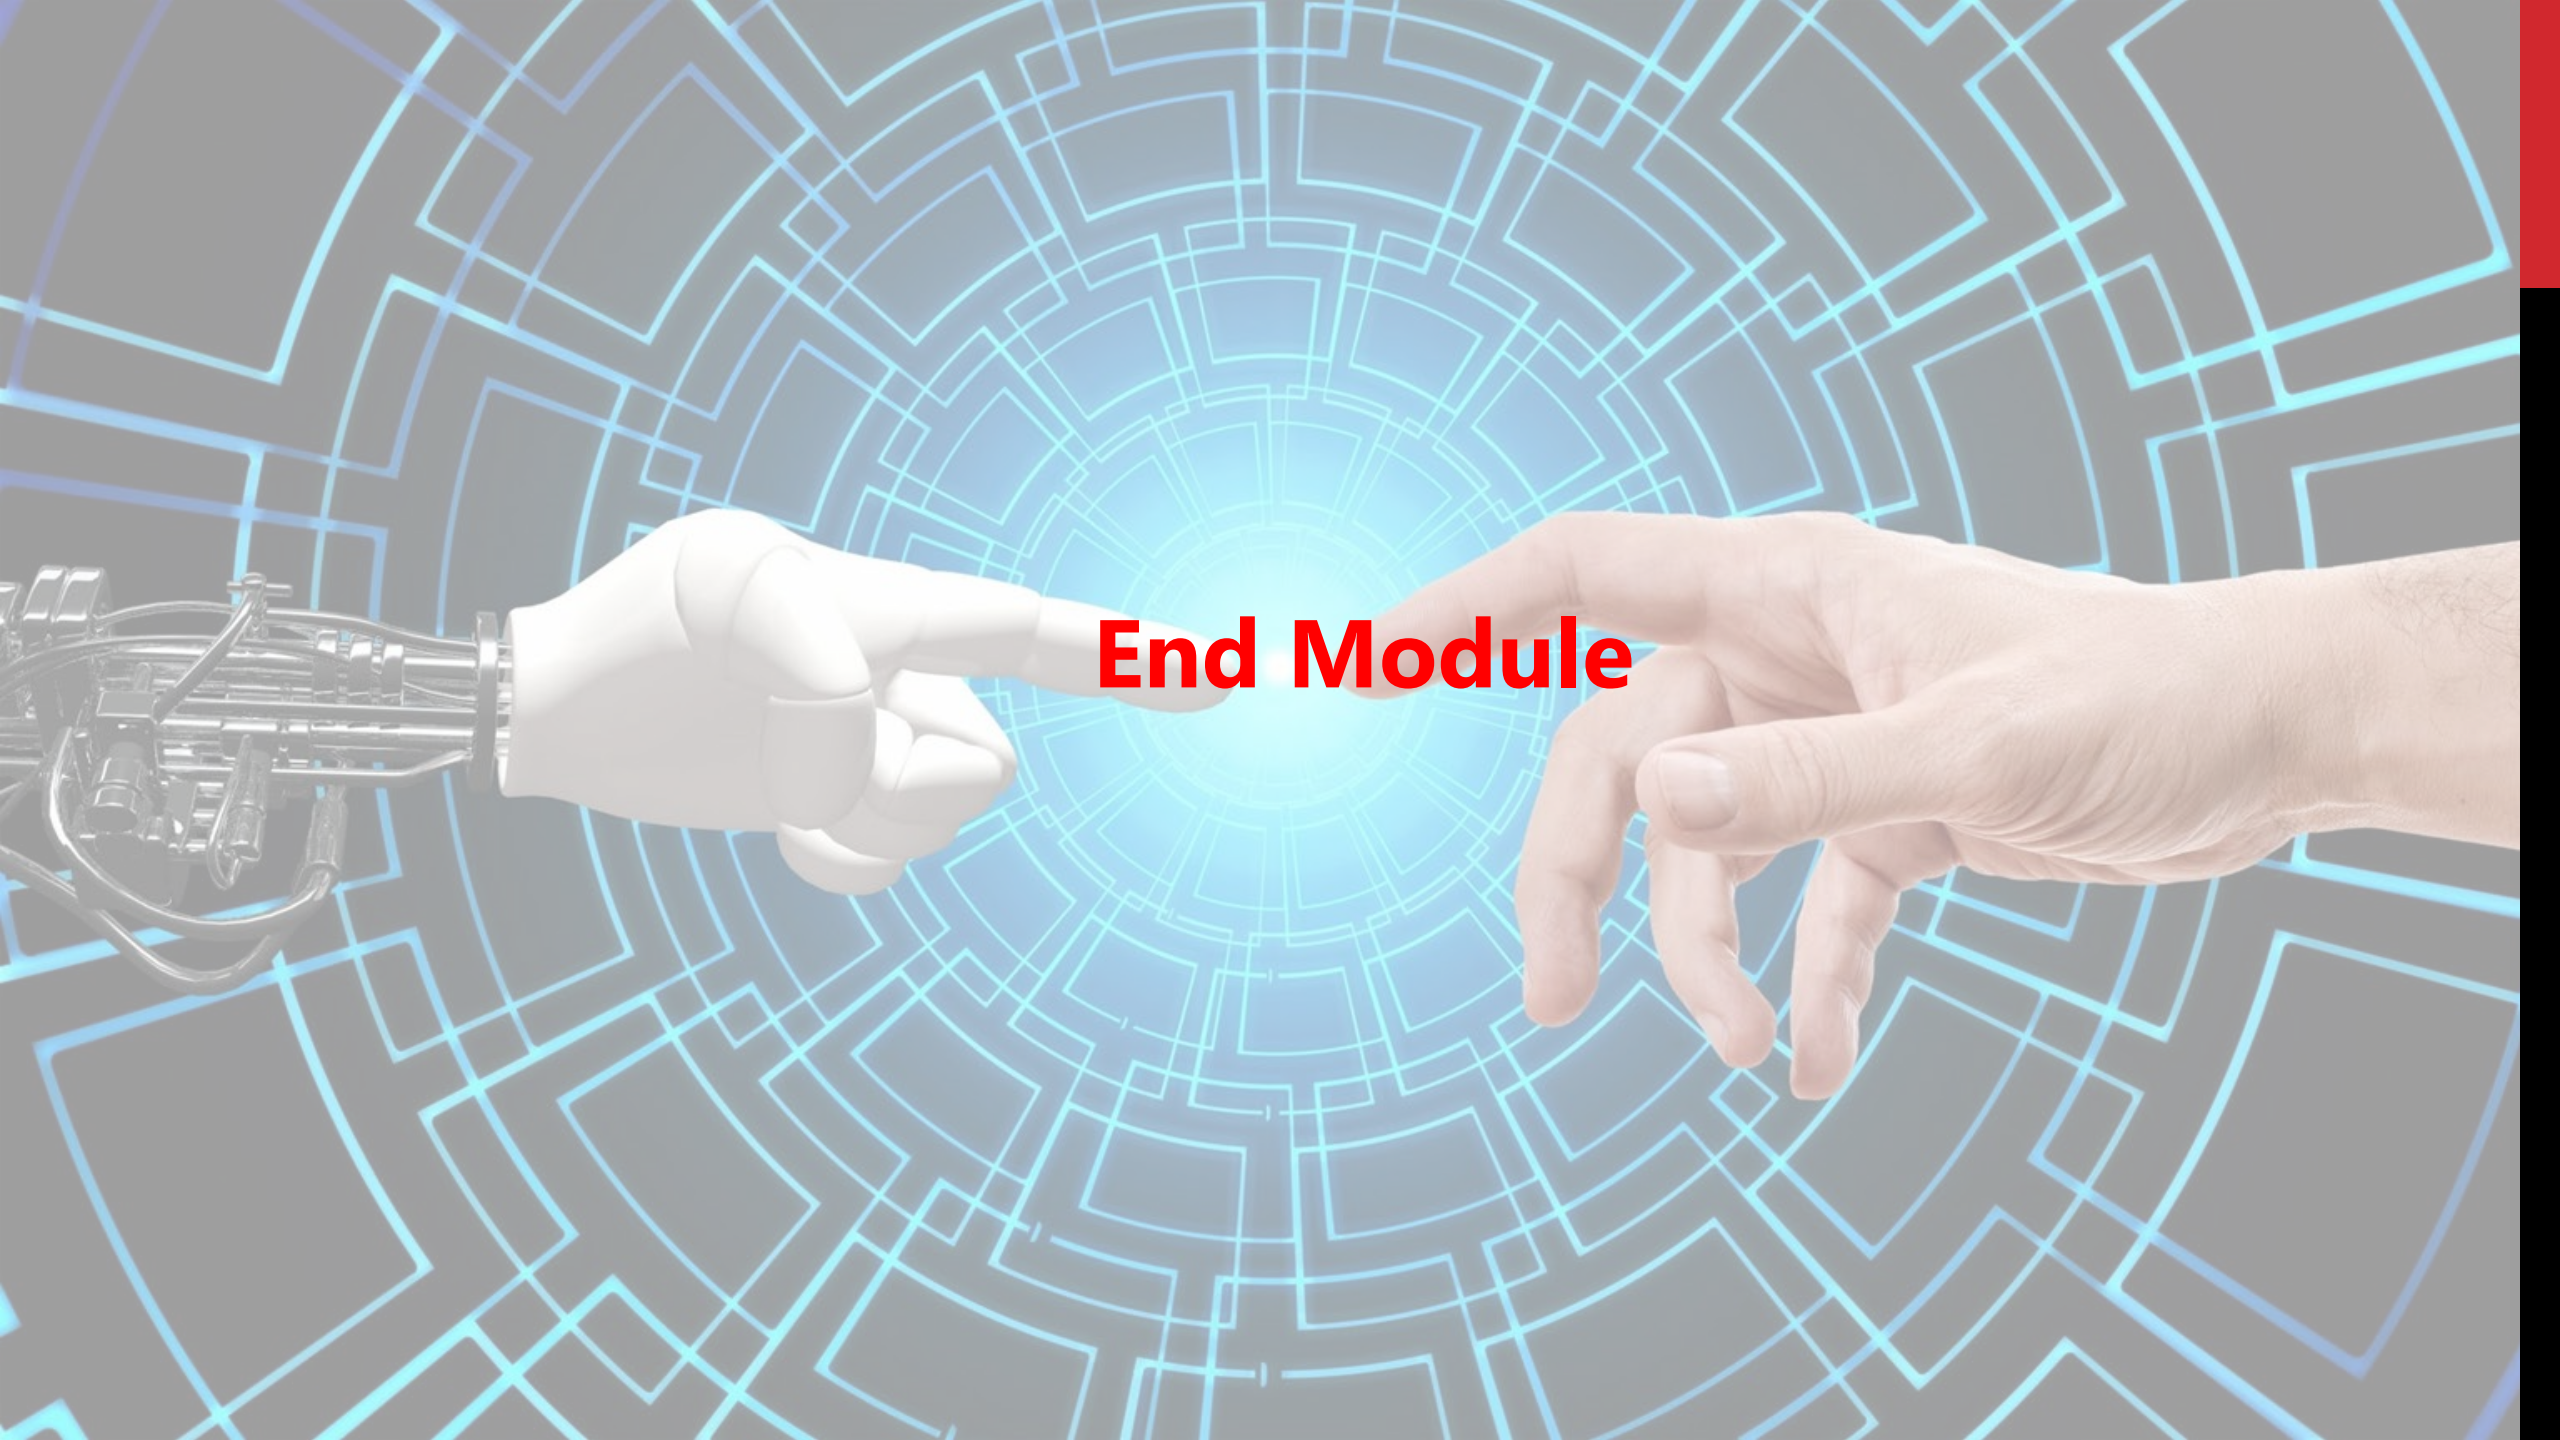


Containers for Managed ML Systems



Build Once, Run Many MLOps Workflow





End Module